

Lexicon Induction for Low-resource Languages

July 18, 2010

Abstract

Statistical machine translation relies on the availability of substantial amounts of human translated texts. Such bilingual resources are available for relatively few language pairs, which presents obstacles to applying current statistical translation models to low-resource languages. In this work, we induce bilingual dictionaries from more plentiful monolingual corpora using a diverse set of cues, including: cross-lingual vector space models, the frequencies of words over time, orthographic similarity, etc. We report the efficacy of these monolingual cues and contrast their performance for a language pair where plentiful bilingual resources are available. We further evaluate the accuracy of bilingual dictionaries induced between English and a set of low resource languages. Since our principal objective is to induce wide coverage lexicons, we contrast the performance of our framework on randomly selected source words with an optimistic results obtained on frequent words and typically reported in lexicon induction literature. Finally, we propose a simple and effective technique for using crowd sourced annotations to incrementally refine the output of our lexicon induction system.

Last sentence doesn't fit too well.

1 Introduction

Statistical methods for machine translation continue to push the state of the art in automatic translation. However, they crucially rely on the availability of large numbers of translations aligned across two languages. Generation of these parallel corpora require the efforts of bilingual speakers and are extremely expensive to produce in sufficient quantities to induce a high quality statistical translation system. As a result, these methods can not be successfully applied to the majority of word's languages and especially those less frequently taught. On the other hand, we now have an unprecedented access to vast and continuously expanding monolingual data. Moreover, these resources often contain cues which could be effectively utilized for inducing bilingual resources, suggesting that we might substantially reduce and eventually eliminate the need for explicitly aligned bilingual translations.

Wordy and vague.

In this work, we exploit these monolingual cues to induce bilingual lexicons between English and a large set of low-density languages. Prior work

on lexicon induction has attempted to automatically learn bilingual lexicons, either by using monolingual corpora ([Rapp, 1999], [Koehn and Knight, 2002], [Schafer and Yarowsky, 2002], [Haghighi et al., 2008] or by exploiting the cross-language evidence of closely related bridge languages that have more resources [Mann and Yarowsky, 2001]. While this work falls into the former category, our objective is substantially different. Unlike the bulk of the previous work, our specific aim is to use monolingual cues to induce large and *broad coverage* dictionaries for a large set of low density languages.

We begin by introducing a set of monolingual cues and studying their efficacy on a language pair for which plentiful bilingual resources are available. We then use them to induce bilingual lexicons between English and a set of ten diverse low resource languages. Finally, we propose a simple and effective technique for using crowd sourced annotations to incrementally refine the output of our lexicon induction system.

2 Related Work

The literature on translation lexicon induction for low-resource languages falls in to two broad categories: 1) utilizing similarity between languages by choosing a high-resource bridge language for translation ([Mann and Yarowsky, 2001], [Schafer and Yarowsky, 2002]) and 2) extracting noisy cues from monolingual corpora and using them to associate words across languages.

In the latter category, [Rapp, 1999] was the first to propose using context of a given word as a clue to its translation. Given a German word with an unknown translation, its surrounding words were collected and translated into English using a small seed dictionary. Words with similar context in a monolingual English corpus were then proposed as translation candidates. The original work employed a relatively large bilingual dictionary containing approximately 16,000 words and tested only on a small collection of 100 manually selected nouns. Subsequent work has explored Rapp’s ideas proposing a variety of alternative similarity metrics, better methods for collecting context, and monolingual cues. [Koehn and Knight, 2002] used a larger test set consisting of the 1000 most frequent words from a German-English lexicon. They also incorporated cues such as frequency and orthographic similarity in addition to context. [Schafer and Yarowsky, 2002] independently proposed using frequency, orthographic similarity and also showed improvements using temporal and word burstiness similarity measures, in addition to context. [Klementiev and Roth, 2006] used phonetic and temporal similarity specifically to discover transliterated terms. [Haghighi et al., 2008] made use of contextual and orthographic cues for learning a generative model from monolingual corpora and a seed lexicon. [Garera et al., 2009] proposed a dependency-based context model that incorporates long-range dependencies, variable context sizes, and reordering. [Mimno et al., 2009], [Boyd-Graber and Blei, 2009] exploited to use multilin-

gual topic models to discover translations.

We first show that evaluation results on a small set of hand selected words typically reported in previous work can be misleading if the objective is to induce large translation dictionaries. We then investigate the effectiveness of individual monolingual cues for inducing broad coverage lexicons for a large number of low-density languages.

3 Inducing Bilingual Lexicons From Monolingual Cues

Various linguistic and corpus cues can be informative for relating word translations across a pair of languages. We begin by considering a few of them and explain how they can be used to measure similarity between a source word and a translation candidate.

*Little technical details,
i.e. no metric defs, etc.
Add?*

3.1 Monolingual cues

Contextual cue. [Rapp, 1999] proposed inducing a translation dictionary from disparate monolingual texts. They populate a bilingual lexicon by projecting *context vectors* across two languages using vector-space semantic models to represent words [Deerwester et al., 1990]. The elements in vector-based representations of a word indicate the frequency of its co-occurrence with all other words in the same language. For instance, the vector representation of “airplane” would indicate that it frequently occurs in contexts near the words “airport”, “flight”, “landing,” “passengers”, “pilot”, “runway”, etc. The similarities of words within one language can be measured using the distance between their vectors, with cosine similarity, for instance. To translate unknown words, [Rapp, 1999] suggests building vector space models of two languages. The elements in an unknown word’s vector are *projected* into the vector space of the other language using the known translations from a small seed bilingual dictionary. This sparse projected vector is compared to the vectors for all words in the target language. The word whose vector is most similar to the projected vector is considered to be the best translation of the unknown word. This process is illustrated in Figure 1.

Temporal cue. Online content is often published along with temporal information: news feeds, for example, are comprised of news stories annotated with date and time of publication. The feeds are specialized for the target geographical locations and vary in content across languages. Still, many events are deemed relevant to multiple audiences and the news stories related to them appear in several languages, although rarely as direct translations of one another. Words associated with these events will appear with increased frequency in multiple languages around the dates when these events are reported. Figure 2 illustrates this idea with temporal histograms of three English words and their Spanish




Figure 1: Lexicon induction using contextual information. First, contextual vectors are projected using a small seed dictionary and then compared with the target language candidates.

translations. Such weak synchronicity provides a cue about the relatedness of words across the two languages, and can be exploited to associate them. In order to score a pair of entities across languages, we can compute the similarity of their temporal signatures.

Orthographic cue. Etymologically related words often retain similar spelling across languages with the same writing system. Capturing these similarities can provide yet another clue about their relatedness. Edit distance defines one such metric, counting the minimal number of edit operations required to transform one string into another. While it won't provide a good signal for most translation pairs, it proves to be a highly effective for related languages such as English and Spanish (see Section 5.1).

Phonetic cue. We can extend the previous idea to language pairs not sharing the same writing system, since many cognates and transliterated words are phonetically similar. [Klementiev and Roth, 2006] train a transliteration model to return high scores to phonetically similar word pairs and show it to be highly successful for Named Entities which are often transliterated. Adding this cue should provide an informative signal for cognates as well.

3.2 Combination Strategies

These cues provide informative and independent means to score a source word and a target candidate, and combining them is likely to produce better lexicons. One idea is to directly combine the ranked lists of induced candidates using the mean reciprocal rank (MRR) heuristic frequently used in Information Retrieval literature. The idea is to score each of the translation candidates with

Figure 2 is from
EuroParl, add the one
from the proposal
instead.




Figure 2: Temporal histograms of three English words and their Spanish translations.

an average reciprocal rank across all rankings, and then sort the candidates in descending order.

The relative informativeness of the cues will depend on the data: e.g. the temporal cue depends on the temporal alignment of both sides of the bilingual corpus, the orthographic cue is uninformative when the two languages use different scripts, etc. However, the MRR heuristic assigns equal weights to each ranking and investigating more suitable combination strategies remains the subject of ongoing work.

Some of the metrics we have introduced (e.g. edit distance) are likely to assign same scores to multiple candidates. We use the following strategy for resolving ties: (1) a set of candidates assigned the same score shares the same rank, (2) a rank assigned to a candidate (a real value) takes into account the number of other candidates with the same or better scores (Figure 3).

3.3 Handling Morphology

Many of the languages in our list such as Russian and Korean are characterized by morphological rules generating a large number of word forms for the same lexeme. Thus, we need to be able to group morphological variants of the same word into an equivalence class and collect their aggregate statistics. For instance, we would like to count the total number of occurrences of $\{Herzegovina,$




Figure 3: Resolving ties.

Hercegovina} on the English side in order to map it accurately to the corresponding equivalence class we may see on the Russian side of our corpus (e.g. {*Герцеговина, Герцеговину, Герцеговины, Герцеговиной*}). In order to keep to our objective of requiring as little language knowledge as possible, we take a rather simplistic approach for grouping equivalence classes. For morphologically rich languages, a set of words are in the same equivalence class if they share a prefix of size five or longer. For the other languages, each unique word is assigned its own class. More sophisticated unsupervised approaches (i.e. [?]) could be incorporated instead, and are a subject of ongoing work.

4 Resources

In this section we briefly describe each of the resources we use for inducing translation lexicons. For those being continuously collected, we give a current snapshot of relevant statistics.

Wikipedia. Wikipedia provides a large repository of continuously edited monolingual texts in a large number of languages. Besides the article content, it also contains various metadata which can be useful for generating additional resources. For example, interlingual links can be used to select articles discussing the same subject across multiple languages. We will use these comparable articles when inducing translations using the contextual cue in Section 5. Table 1 lists counts of such page pairs between English and 42 other languages.

Dictionaries and stopword lists. Dictionaries provide a crucial resource for projecting vector-base representations of equivalence classes and for evaluating the quality of induced lexicons. Table 1 lists the sizes of our dictionaries, along with their coverage of the English side of the “comparable” bilingual Wikipedia subsets in column 2. Type coverage measures the proportion of unique Wikipedia English words which have an entry in the dictionary, while token coverage takes into account their corpus counts. Low type coverage for most languages is primarily due to misspelled or incorrectly extracted article text in Wikipedia

Language	Wikipedia page pairs	Dictionary		
		entries	src type/token coverage (%)	script
Tigrinya	36	56	0.3/0.6	r
Punjabi	401	76,311	12.2/57.7	r/o
Kyrgyz	492	74,890	2.6/42.4	r
Somali	585	230	0.2/6.8	r
Nepali	1,293	6,812	3.6/39.7	r
Tibetan	1,358	59,083	2.5/35.1	r
Uighur	1,814	16,285	1.7/37.9	r
Maltese	1,896	7,574	2.9/44.7	r
Turkmen	3,137	91,928	4.6/54	r
Kazakh	3,470	145,750	1/34	r
Mongolian	4,009	948	0.3/20.7	r
Tatar	4,180	8,557	1.7/33	c/r
Kurdish	5,059	9,870	1.2/32	r
Uzbek	5,875	190,688	6.2/66.4	r
Kapampangan	6,827	1,000	0.2/7	r
Urdu	7,674	36,428	2.9/57.1	r
Irish	9,859	887	0.2/23.5	r
Azeri	12,568	231,891	0.6/42.7	r
Tamil	13,470	165,004	2.6/43.6	o
Albanian	13,714	188,563	5.4/55.6	r
Afrikaans	14,315	11,389	0.8/50.2	r
Hindi	14,824	58,179	2.2/37.7	r
Bangla	16,026	1,606	0.2/20.6	o/r
Tagalog	17,757	247,662	3.5/65.8	r
Latvian	22,737	148,363	3.9/65.5	r
Bosnian	23,144	18,283	1/44.5	r
Welsh	25,292	25,832	2/50.1	r
Latin	31,195	18,884	0.9/36.3	r
Basque	38,594	880	0.1/7	r
Thai	40,182	14,925	0.5/36.4	o
Farsi	58,651	198,605	2.6/60.5	a/r
Bulgarian	68,446	316,631	3/66.3	c/r
Serbian	71,018	168,140	3.7/71.2	r
Indonesian	73,962	67,633	1.1/61.5	r
Slovak	76,421	233,093	2.6/65	r
Korean	84,385	229,742	2.2/66.5	k
Turkish	86,277	1,272,881	3.3/47.9	r
Ukrainian	91,022	14,056	0.5/35.3	r
Romanian	97,351	249,479	2.4/64.3	r
Russian	295,944	423,009	1.6/57.5	c
Spanish	371,130	347,441	1.7/54.7	r
Polish	438,053	261,463	1.3/57.8	r

Table 1: Number of Wikipedia pages that have interlanguage links with English, along with dictionary statistics. The last column contains the dictionary script: roman (r), cyrillic (c), arabic (a), korean (k), or other (o).

articles generating a long tail of low frequency words, and noise in the dictionaries themselves. On the target side, dictionary noise and rich morphology of some of the languages we consider also substantially reduces coverage. Moreover, some of the dictionaries contain translations whose script was changed to

Latin (see the last column of Table 1). Due to these reasons, evaluation results we present below tend to be conservative. We propose an alternative method of obtaining small translation lexicons in Section 5.4. Finally, we have also collected stop words for Farsi, Bangla, Hindi, Polish, Spanish, Russian, Romanian, and English.

News wire. We have assembled additional monolingual corpora by continuously crawling a set of regularly updated news websites. Up to 10 years of collected and processed data are listed in Table 2 for 25 languages: the last page counts a number of tokens collected and the previous columns list the number of days with at least one associated story for each of the 10 years. We used a set of heuristics to extract language and temporal information from page URL and body text; a substantial portion of the crawled data not included in the table still remains to be processed for language and time.

Language	Pages per year											Total tokens
	'00	'01	'02	'03	'04	'05	'06	'07	'08	'09	'10	
Kyrgyz	-	-	-	-	-	-	1	7	5	15	40	126,217
Somali	-	-	-	-	-	-	-	-	28	113	64	421,589
Nepali	-	-	-	-	1	3	4	2	6	13	8	59,309
Kazakh	-	-	-	-	-	-	-	-	1	45	32	1,839,748
Uzbek	-	-	-	-	-	-	22	29	79	49	33	507,320
Urdu	-	-	1	44	33	8	36	30	76	281	98	2,874,969
Tamil	-	-	-	-	-	1	1	1	13	28	27	61,990
Albanian	1	-	-	-	2	26	13	98	94	120	70	828,005
Hindi	-	14	181	285	283	332	334	325	331	349	102	12,589,950
Bangla	-	-	-	-	-	1	1	-	4	18	29	129,984
Latvian	-	1	1	1	1	1	3	21	178	341	75	28,309,813
Bosnian	-	-	-	-	-	-	1	-	-	-	41	79,206
Welsh	-	2	-	142	245	129	1	1	1	5	-	2,603,551
Farsi	-	10	49	114	160	301	297	339	366	365	109	26,041,987
Serbian	-	1	1	-	-	-	3	-	14	84	19	221,503
Indonesian	-	-	-	-	-	-	-	1	1	59	99	1,135,783
Slovak	1	-	-	-	1	1	-	158	356	364	100	103,732,925
Turkish	-	-	-	1	7	12	8	9	24	153	69	1,135,200
Ukrainan	-	-	-	15	26	30	70	89	164	102	76	1,254,852
Russian	46	365	363	353	350	353	353	358	365	365	100	47,857,954
Spanish	-	313	352	364	366	365	365	365	366	365	104	59,732,042
English	366	365	365	365	366	365	365	365	366	365	187	1,090,171,115
Arabic	-	-	-	-	-	-	-	-	-	144	100	1,189,680
Pashto	-	-	-	-	-	3	10	9	7	26	37	520,450
Chinese	-	-	-	-	-	5	1	6	7	14	32	1,864,565

Table 2: Crawled newswire pages with identified language and publication date information.

Note: Ben's runs used an older crawl from April.

Parallel texts. EuroParl [Koehn, 2005] is a parallel corpus in 11 languages compiled from European parliament proceedings published on the web. Besides sentence alignment, a portion of the corpus contains temporal annotation. While we cannot assume such resources to be available for all of our language pairs, we use it to get a sense of an upper bound of the performance we can achieve

with our methods.

5 Experimental Evaluation

We evaluate performance of the lexicon induction framework on the monolingual resources we have collected and described in Section 4. In the first set of experiments, we consider a high resource language pair to establish relative efficacy of the cues and to get a sense of an upper bound on the overall performance. We then induce translations between English and ten low resource languages. Since many of the seed and test dictionaries are sparse and noisy, we investigate the feasibility of crowd-sourcing annotations and their use in an iterative induction procedure.

For each language, we construct a test dictionary by randomly selecting 10% of the entries in the corresponding bilingual dictionary (see Section 4), and leave the remainder as a seed lexicon for the contextual cue. We select 1000 source words from the test dictionary and compute the accuracy of the induced translations. The performance is measured by top- k accuracy, i.e. the proportion of the source words which have at least one test dictionary translation among top k of its induced candidates. Note that the reported performance is conservative since the test dictionaries are both noisy and sparse. Unless mentioned otherwise, we use the simple equivalence class heuristic (i.e. each unique token is assigned its own equivalence class) when constructing both contextual vectors and candidates. We do not consider stop words as both source and target candidates; for languages without a stop word list, we remove 200 most frequent equivalence classes.

As we have pointed out, one of the objectives of this work is to induce lexicons with wide coverage. Thus, unlike most of the previous work, we are particularly interested in how well we can do on inducing translations for source equivalence classes selected *randomly* from the test dictionary.

5.1 Lexicon induction for a high-resource language pair

We begin by evaluating the relative performance of lexicons induced from contextual, temporal, and orthographic cues as well as the MRR rank aggregation scheme on the English-Spanish EuroParl parallel data for inducing translations for most frequent and random 1000 words source present in our test dictionary (top left and right of Figure 4, respectively). The corpus is time stamped (spanning 656 days) and perfectly *temporally* aligned, so it is not surprising that the temporal cue provides a strong signal for inducing translations. All of these cues are informative and orthogonal, so combining them substantially improves results with 73% and 82% accuracy at top 100 and 500 for most frequent source words, and 70% and 79%, respectively, for random source tokens.




Figure 4: Accuracy on parallel en-es data for most frequent (top left) and random (top right) 1000 source words, and on wikipedia / newswire for most frequent (bottom left) and random (bottom right) 1000 source words.

We repeat the same experiment using wikipedia and newswire data to derive dictionaries from contextual and temporal cues, respectively (Figure 4, bottom). Since the newswire corpora is only weakly temporally aligned, the performance of the temporal cue drops substantially. Still, the cues provide informative and non-redundant signals, which can be combined to obtain better quality lexicons. Notice the substantial drop in performance (about **19%**) for randomly selected source tokens. While not entirely surprising, it supports our observation that evaluating on frequent source words alone can be misleading if the objective is to induce wide coverage dictionaries.

5.2 Context vs. alignments

Word alignments can be induced directly from sentence aligned corpora. In this set of experiments, we compare bilingual lexicons derived from word alignments (produced by GIZA++) to those generated from the contextual cue alone on the English-Spanish section of the EuroParl corpus (see Figure 5). While word

alignments induce a much more informative signal than context alone, sufficient sentence aligned bilingual data is not available for most of the low resource languages we consider in this work.




Figure 5: Accuracy of lexicons induced from alignments and context on parallel en-es data for random and most frequent 1000 source words.

5.3 Lexicon induction for a low-resource languages

Next, we selected 10 low resource languages for which we have collected sufficient amount of newswire data and have seed dictionaries in an appropriate script: Welsh, Farsi, Indonesian, Latvian, Russian, Slovak, Albanian, Serbian, Turkish, and Uzbek. Figure 6 shows the accuracy of lexicons induced from the contextual (top), and temporal (middle) cues, and the MRR aggregation scheme (bottom); results for 1000 most frequent and random source words are shown on the left and right, respectively.

Lexicon accuracy drops substantially for Russian and Turkish on random source words, as expected. For the remaining languages, low dictionary coverage means a large overlap between 1000 random and most frequent test dictionary terms. That is, both sets of experiments effectively test the accuracy on random source words, since test dictionary entries themselves were selected at random. Both temporal and contextual cues are informative, but the simple aggregation heuristic fails to provide an improvement. While simple ensemble schemes typically work well for a large number of diverse signals, a more sophisticated approach to combining few available ranked lists remains the subject of our ongoing work.

Confusing? Drop most frequent?

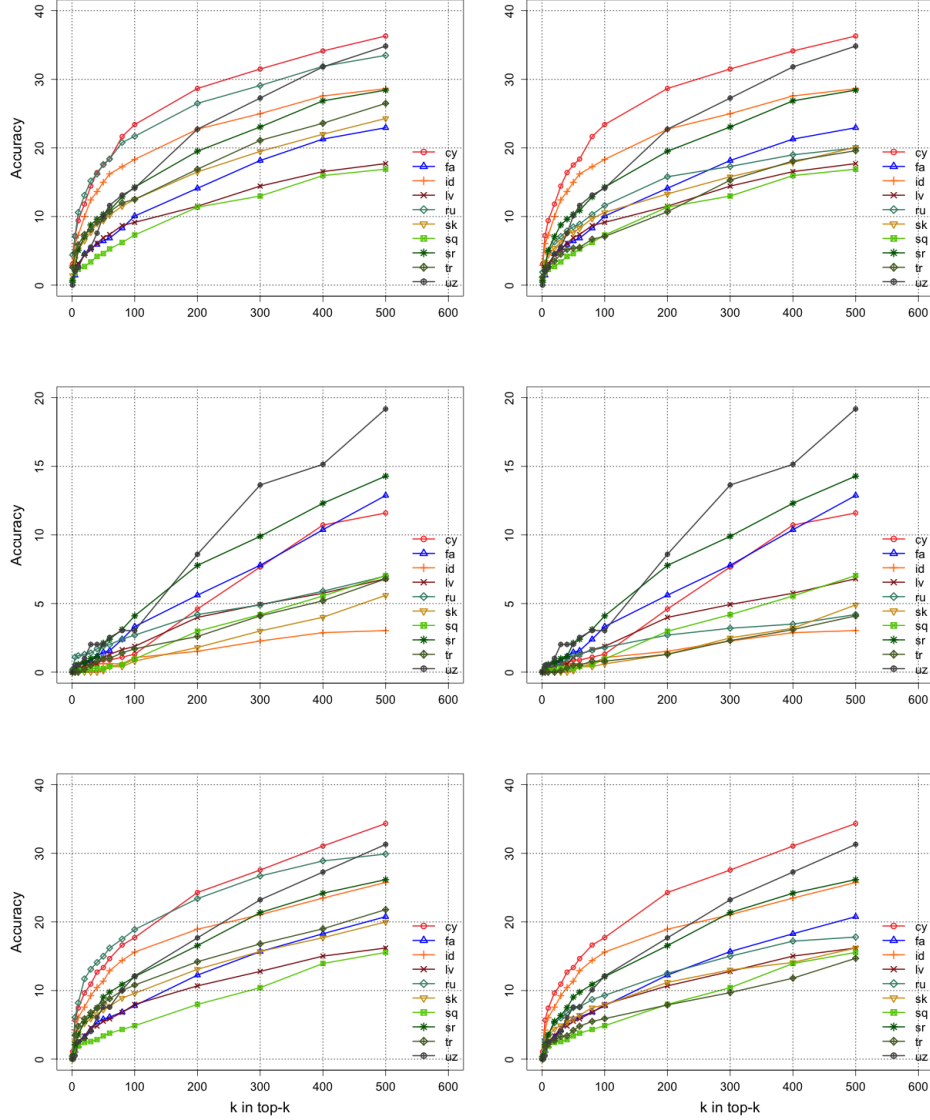


Figure 6: Accuracy of contextual (top) and temporal (middle) cues and the MRR scheme (bottom) for most frequent (left) and random (right) 1000 source words for Welsh (cy), Farsi (fa), Indonesian (id), Latvian (lv), Russian (ru), Slovak (sk), Albanian (sq), Serbian (sr), Turkish (tr), and Uzbek (uz).

5.4 Crowd-sourcing annotation

We use our existing bilingual dictionaries to induce large bilingual lexicons via the contextual cue and to evaluate their accuracy. However, these dic-

tionaries vary substantially in quality and coverage across languages and corpora (see Section 4). In this set of experiments we follow previous work on crowd-sourcing annotations [Snow et al., 2008, Callison-Burch, 2009] and study [Irvine and Klementiev, 2010] the viability of obtaining translations for a low-resource languages specifically for use in our induction framework. First, we use the contextual cue to induce lexical translation pairs from the Wikipedia monolingual data. Then, we pay Amazon Mechanical Turk (MTurk) workers a small amount to check and correct our system output. We can then use the updated lexicons to inform another iteration of lexicon induction, gather a second set of MTurk annotations, and so on.

For 32 of the 42 languages in Table 1, we were able to induce lexical translation candidates and post them on MTurk for annotation. For these languages we presented annotators with top ten scored candidates for a set of 100 English words and asked them to mark correct translations. If our seed dictionary included an entry for a source word, we included the translation in the candidate list as a positive control. Additionally, we included a random word in the foreign language as a negative control. We do not have dictionaries for the remaining 10 languages, so we asked workers to type translations for 100 English words. We had three distinct workers provide such annotations for each source word.

Figure 7 (top) shows the time it took to complete annotation of for 37 languages on MTurk. Annotations for the following languages were posted for a week and were never completed: Tigrinya, Uighur, Tibetan, Kyrgyz, and Kazakh. All five of the uncompleted required typing annotations, a more time consuming task than checking translation candidates. Not surprisingly, languages with many speakers (Hindi, Spanish, and Russian) and languages spoken in and near India (Hindi, Tamil, Urdu) were completed very quickly. Figure 7 (bottom) shows the percent of positive control candidate translations that were checked by the majority of workers (at least two of three). The highest amounts of agreement with the controls were for Spanish and Polish, which indicates that those workers completed the annotations more accurately than the workers who completed, for example, the Tatar and Thai annotations. However, the seed dictionaries are very noisy, so this finding may be confounded by discrepancies in the quality of our dictionaries. The noisy dictionaries also explain why agreement with the positive controls is, in general, relatively low.

To understand the utility of MTurk generated translation for inducing lexicons, we supplemented our dictionaries for each of the 37 languages for which we gathered MTurk annotations with translation pairs that workers agreed were good (both chosen from the candidate set and manually translated). We compared seed dictionaries of size 200 with those supplemented with, on average, 69 translation pairs. We found an average relative increase in accuracy of our output candidate set (evaluated against complete available dictionaries) of **53%**.

In sum, we found that the iterative approach of automatically generating




Figure 7: Top: time to complete annotation of 100 English words. Division of the time between posting and the completion of the first annotation unit (HIT) and the time between the completion of the first and last HIT shown. HITs that required lexical translation only are marked with an *. Bottom: percent of positive control candidate translations for which two or three workers checked as accurate.

noisy annotation and asking MTurk users to correct it to be an effective means of obtaining supervision. These correction tasks are simple, can be completed quickly for a large number of low resource languages, and produce high quality annotation.

6 System Overview

In this section we touch on some of the implementation details of the lexicon induction framework: we overview the data collection and lexicon induction procedures and explain how the framework can be extended to include new monolingual resources and cues derived from them.

6.1 Data Collection

While some monolingual resources (see Section 4) are static, others require ongoing collection. We have set up the nutch crawler¹ to continuously crawl a number of web sites generating news content in the languages of interest. The crawl results are periodically processed (see Figure 8) to (1) parse page content and extract metadata associated with the page, (2) merge with previously extracted pages, (3) identify language of the page content, and (4) generate time annotated corpora for each of the languages. See Table 1 for a current summary of the collected data.




Figure 8: Ongoing data collection.

6.2 Lexicon Induction Procedure

Let us turn to the induction procedure, highlight some of the most relevant code, and show how the framework can be extended to include new monolingual resources and cues. Figure 9 shows the implementation layout and Figure 10 gives a high level view of the lexicon induction procedure.

We argued in Section 3.3 that collecting aggregate statistics for morphological variants of a single lexeme is important when dealing with morphologically rich languages. Base class `EquivalenceClass` groups morphological variants present in the data into equivalence classes and maintains a set of aggregate statistics derived from monolingual cues. In turn, each of the statistics, or properties, is implemented by a subclass of `Property`.

The induction procedure begins with two passes through both source and target language monolingual corpora (step 1 and 2 on Figure 10, respectively) implemented in `DataPreparer`. Each of the available corpora (e.g. see Section 4) is accessed through a corresponding subclass of `CorpusAccessor`. In the first pass, morphological variants are collected (see `EquivalenceClassCollector`) to generate the corresponding equivalence classes and, in the second pass, a set of properties such as contextual vectors, temporal distributions, topic information, etc. is collected for each of the equivalence classes. The initial set of equivalence classes is pruned by a series of filters extending

¹<http://nutch.apache.org/>

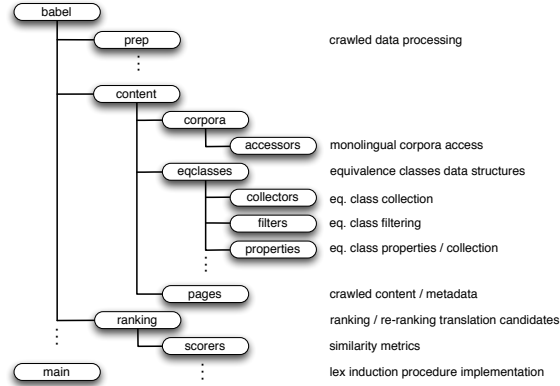


Figure 9: Package layout of the lexicon induction framework.

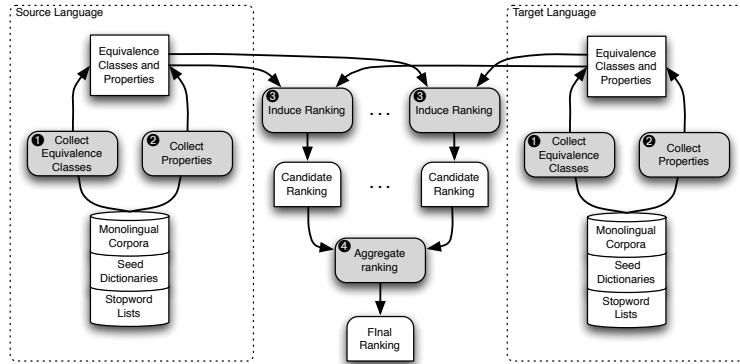


Figure 10: A high level overview of the lexicon induction framework. Equivalence classes and corresponding properties are first extracted from monolingual data (steps 1 and 2). Similarity metrics defined over the properties are then used to produce rankings over target candidates in step 3. Finally, ranked lists are aggregated to produce the final rankings in step 4.

EquivalenceClassFilter in order to throw out patently incorrect or undesirable candidates, e.g. stop words, least or most frequent classes, strings containing numbers or letters of a wrong script, etc. Both source and target equivalence classes along with the collected statistics are persisted on disk.

Next, collected properties along with the corresponding similarity metrics extending **Scorer** are used to produce a ranked list of candidates for each of the source equivalence classes. This step involves a substantial amount of computation since each of the source equivalence classes is compared with all of the target candidates. Its implementation in **Ranker** is parallelized, which substantially speeds up this step.

Ranked candidate lists induced for each source equivalence class from multiple cues are aggregated (see **Reranker**) into a joint ranking in step 4 on Figure 10. Finally, the induced ranked candidate lists are evaluated in **NBestCollector**.

Listing 1 shows an example configuration file for setting up the induction process. It is split into 5 sections, one would:

- The **corpora** section lists both source and target monolingual corpora with additional configuration parameters specific to the corresponding subclass of **CorpusAccessor**.
- The **resources** section specifies additional resources, such as stop word lists and bilingual dictionaries.
- The **preprocessing** section configures the two stage preprocessing stage, i.e. which resources to use to generate equivalence classes and how to collect their properties. For example, the **candidates** section on Listing 1 specifies that the simple and prefix heuristics (see Section 3.3) should be used for generating source and target equivalence classes, respectively, and that the classes should be pruned if they occur fewer than 10 times in the data.
- Finally, the **experiments** section configures the induction process. The configuration parameters can be used to choose most frequent or random source equivalence classes for induction (**RandomSource**), the portion of the dictionary to use for projecting contextual vectors (**DictionaryPercentToUse**), the target candidate ranked lists size to induce (**NumTranslationsToAddPerSource**), the number of threads to use when generating rankings (**NumRankingThreads**), and to specify which properties are to be used to induce those rankings and whether or not to aggregate them (**DoAggregate**).

In order to extend the framework to add a new monolingual resource and/or include additional cues:

- Extend **CorpusAccessor** to enable access to a new resource.
- Extend **Property** and **PropertyCollector** to manage and collect desired statistics from a monolingual resource.
- Extend **Scorer** to implement a similarity metric for scoring a source and a target candidate equivalence classes.

```

<?xml version="1.0" encoding="UTF-8" standalone="no"?>
<configuration>
  <corpora>
    <wiki>
      <Path>./resources/wiki/en-ru</Path>
      <SrcRegExp>.*\\.en</SrcRegExp>
      <TrgRegExp>.*\\.ru</TrgRegExp>
    </wiki>
    <crawls>
      <Path>./resources/crawls</Path>
      <SrcSubDir>en</SrcSubDir>
      <TrgSubDir>ru</TrgSubDir>
      <DateFrom>00-01-01</DateFrom>
      <DateTo>10-04-20</DateTo>
    </crawls>
    <europarl>
      <Path>./resources/es-en</Path>
      <SrcSubDir>en</SrcSubDir>
      <TrgSubDir>es</TrgSubDir>
      <DateFrom>96-04-15</DateFrom>
      <DateTo>06-10-13</DateTo>
    </europarl>
  </corpora>
  <resources>
    <stopwords>
      <Path>./resources/stopwords</Path>
      <SrcStopWords>en.stop</SrcStopWords>
      <TrgStopWords>ru.stop</TrgStopWords>
    </stopwords>
    <dictionary>
      <Path>./resources/dictionaries</Path>
      <Dictionary>en-ru.dict</Dictionary>
    </dictionary>
  </resources>
  <preprocessing>
    <Path>./preprocessing</Path>
    <FilterRomanTrg>false</FilterRomanTrg>
    <input>
      <Context>wiki</Context>
      <Time>crawls</Time>
    </input>
    <candidates>
      <SrcEqClass>babel.content.eqclasses.SimpleEquivalenceClass</SrcEqClass>
      <TrgEqClass>babel.content.eqclasses.PrefixEquivalenceClass</TrgEqClass>
      <PruneIfOccursMoreThan>-1</PruneIfOccursMoreThan>
      <PruneIfOccursFewerThan>10</PruneIfOccursFewerThan>
      <PruneMostFrequentSrc>-1</PruneMostFrequentSrc>
      <PruneMostFrequentTrg>-1</PruneMostFrequentTrg>
    </candidates>
    <context>
      <SrcEqClass>babel.content.eqclasses.SimpleEquivalenceClass</SrcEqClass>
      <TrgEqClass>babel.content.eqclasses.SimpleEquivalenceClass</TrgEqClass>
      <PruneEqIfOccursMoreThan>-1</PruneEqIfOccursMoreThan>
      <PruneEqIfOccursFewerThan>5</PruneEqIfOccursFewerThan>
      <PruneContextToSize>-1</PruneContextToSize>
      <Window>2</Window>
    </context>
    <time>
      <Align>true</Align>
    </time>
  </preprocessing>
  <output>
    <Path>./output</Path>
  </output>
  <experiments>
    <time>
      <SlidingWindow>false</SlidingWindow>
      <WindowSize>1</WindowSize>
    </time>
    <RandomSource>false</RandomSource>
    <NumSource>1000</NumSource>
    <NumTranslationsToAddPerSource>500</NumTranslationsToAddPerSource>
    <DictionaryPercentToUse>0.9</DictionaryPercentToUse>
    <DictionaryPruneNumTranslations>-1</DictionaryPruneNumTranslations>
    <NumRankingThreads>15</NumRankingThreads>
    <DoTime>true</DoTime>
    <DoContext>true</DoContext>
    <DoEditDistance>false</DoEditDistance>
    <DoAggregate>true</DoAggregate></experiments>
  </configuration>

```

Listing 1: Example configuration file.

7 Conclusions and Future Work

- Investigated the efficacy of monolingual cues for a number of low-density languages / resources. If the objective is wide coverage, results reported in prior work are misleading.
- Bilingual lexicons is preliminary step for low resource MT. Incorporate metric scores as features.
- Aggregation schemes for few signals and some amount of supervised data.

References

- [Boyd-Graber and Blei, 2009] Boyd-Graber, J. and Blei, D. M. (2009). Multilingual topic models for unaligned text. In *Proc. of the Conference on Uncertainty in Artificial Intelligence (UAI)*.
- [Callison-Burch, 2009] Callison-Burch, C. (2009). Fast, cheap, and creative: Evaluating translation quality using amazons mechanical turk. In *Proc. of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- [Deerwester et al., 1990] Deerwester, S., Dumais, S. T., Furnas, G. W., Landauer, T. K., and Harshman, R. (1990). Indexing by latent semantic analysis. *Journal of the American Society for Information Science*, 41(6):391–407.
- [Garera et al., 2009] Garera, N., Callison-Burch, C., and Yarowsky, D. (2009). Improving translation lexicon induction from monolingual corpora via dependency contexts and part-of-speech equivalences. In *Proc. of the Annual Conference on Computational Natural Language Learning (CoNLL)*.
- [Haghighi et al., 2008] Haghighi, A., Liang, P., Berg-Kirkpatrick, T., and Klein, D. (2008). Learning bilingual lexicons from monolingual corpora. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*.
- [Irvine and Klementiev, 2010] Irvine, A. and Klementiev, A. (2010). Using mechanical turk to annotate lexicons for less commonly used languages. In *The NAACL Workshop on Creating Speech and Language Data With Amazon’s Mechanical Turk*.
- [Klementiev and Roth, 2006] Klementiev, A. and Roth, D. (2006). Weakly supervised named entity transliteration and discovery from multilingual comparable corpora. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*.
- [Koehn, 2005] Koehn, P. (2005). Europarl: A parallel corpus for statistical machine translation. In *Proc. of the Machine Translation Summit*.
- [Koehn and Knight, 2002] Koehn, P. and Knight, K. (2002). Learning a translation lexicon from monolingual corpora. In *ACL Workshop on Unsupervised Lexical Acquisition*.
- [Mann and Yarowsky, 2001] Mann, G. S. and Yarowsky, D. (2001). Multipath translation lexicon induction via bridge languages. In *Proc. of the Annual Meeting of the North American Association of Computational Linguistics (NAACL)*.
- [Mimno et al., 2009] Mimno, D., Wallach, H., Naradowsky, J., Smith, D., and McCallum, A. (2009). Polylingual topic models. In *Proc. of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

- [Rapp, 1999] Rapp, R. (1999). Automatic identification of word translations from unrelated english and german corpora. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 519–526.
- [Schafer and Yarowsky, 2002] Schafer, C. and Yarowsky, D. (2002). Inducing translation lexicons via diverse similarity measures and bridge languages. In *Proc. of the Annual Conference on Computational Natural Language Learning (CoNLL)*, pages 146–152.
- [Snow et al., 2008] Snow, R., OConnor, B., Jurafsky, D., and Ng, A. Y. (2008). Cheap and fast - but is it good? evaluating non-expert annotations for natural language tasks. In *Proc. of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*.